

# Assg1\_1

April 28, 2019

```
In [77]: %matplotlib inline
```

```
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV

from sklearn.metrics import accuracy_score
from sklearn import metrics as m
from sklearn.model_selection import StratifiedShuffleSplit
from sklearn.metrics import roc_auc_score

from sklearn.svm import LinearSVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.neural_network import MLPClassifier

from sklearn.metrics import roc_curve
from sklearn.metrics import auc
from sklearn.metrics import classification_report, confusion_matrix

import tqdm as tqdm
```

```
In [78]: import warnings
warnings.filterwarnings("ignore")
```

## Helper functions

```
In [79]: # decode categories from encoded values
def getCategories(attribute):
    print("{} categories".format(attribute))
    categories = df[attribute].astype('category').cat.categories
    for x, y in zip(categories, [l for l in range(1, len(categories) + 1)] ):
        print(y, x)

def roc(classifiers, X_train, y_train, X_test, y_test):
```

```

for name, classifier in classifiers.items():
    cls = classifier
    cls = cls.fit(X_train, y_train)

    y_pred = cls.predict(X_test)

    y_true = sorted(y_test)
    y_score = sorted(y_pred)
    # Compute fpr, tpr, thresholds and roc auc
    fpr, tpr, thresholds = roc_curve(y_true, y_score)
    roc_auc = auc(fpr, tpr)

    # Plot ROC curve
    plt.plot(fpr, tpr, label='ROC curve (area = %0.3f)' % roc_auc)
    plt.plot([0, 1], [0, 1], 'k--') # random predictions curve
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.0])
    plt.xlabel('False Positive Rate or (1 - Specificity)')
    plt.ylabel('True Positive Rate or (Sensitivity)')
    plt.title('Receiver Operating Characteristic')
    plt.legend(loc="lower right")

def classify_and_compare(classifiers, X_train, y_train, X_test, y_test):
    metric_columns = ["Classifier", "Accuracy", "Precision Score", "Recall Score", "F1-Score"]
    df_metric = pd.DataFrame(columns=metric_columns)

    for name, classifier in classifiers.items():
        cls = classifier
        cls = cls.fit(X_train, y_train)

        y_pred = cls.predict(X_test)

        accuracy = m.accuracy_score(y_test, y_pred)
        precision = m.precision_score(y_test, y_pred, average='macro')
        recall = m.recall_score(y_test, y_pred, average='macro')
        roc_auc = roc_auc_score(y_test, y_pred)
        f1_score = m.f1_score(y_test, y_pred, average='macro')

        metric_entry = pd.DataFrame([[name, accuracy, precision, recall, f1_score, roc_auc]])
        df_metric = df_metric.append(metric_entry)

    print(name)
    print(classification_report(y_test, y_pred))

print(df_metric)
plt.xlabel('Accuracy')
plt.title('Classifier Accuracy')

```

```

sns.set_color_codes("muted")
sns.barplot(x='Accuracy', y='Classifier', data=df_metric, color="g")
plt.show()

def findBestParamters(clf, parameters, X_train, y_train):
    clf =GridSearchCV(clf, parameters)
    clf.fit(X_train, y_train)

    print("Best parameters are: {}".format(clf.best_params_))

In [80]: df = pd.read_csv('./dataset/[UCI]bank-additional/bank-additional/bank-additional-full

df_x = df.copy()
del df_x['y']

df_y = df[['y']].copy().astype('category')
# label encode y - yes->1 no->0
df_y = df_y.apply(lambda x: x.cat.codes)

df.head()

Out[80]:
   age  job  marital  education  default  housing  loan  contact  \
0   56  housemaid  married  basic.4y      no      no   no  telephone
1   57  services  married  high.school  unknown      no   no  telephone
2   37  services  married  high.school      no    yes   no  telephone
3   40   admin.  married  basic.6y      no      no   no  telephone
4   56  services  married  high.school      no      no  yes  telephone

   month  day_of_week  ...  campaign  pdays  previous  poutcome  emp.var.rate  \
0   may             mon  ...         1    999         0  nonexistent         1.1
1   may             mon  ...         1    999         0  nonexistent         1.1
2   may             mon  ...         1    999         0  nonexistent         1.1
3   may             mon  ...         1    999         0  nonexistent         1.1
4   may             mon  ...         1    999         0  nonexistent         1.1

   cons.price.idx  cons.conf.idx  euribor3m  nr.employed  y
0          93.994         -36.4      4.857      5191.0  no
1          93.994         -36.4      4.857      5191.0  no
2          93.994         -36.4      4.857      5191.0  no
3          93.994         -36.4      4.857      5191.0  no
4          93.994         -36.4      4.857      5191.0  no

[5 rows x 21 columns]

```

---

## 1 Analysis

---

```
In [81]: # shape of the feature dataframe
print(df_x.shape)
```

```
(41188, 20)
```

```
In [82]: # features
print(list(df_x.columns))
```

```
['age', 'job', 'marital', 'education', 'default', 'housing', 'loan', 'contact', 'month', 'day_of_week', 'duration', 'campaign', 'pdays', 'previous', 'poutcome', 'emp.var.rate', 'cons.price.idx', 'cons.conf.idx', 'euribor3m', 'nr.employed']
```

```
In [83]: print(df_x.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 41188 entries, 0 to 41187
Data columns (total 20 columns):
age                41188 non-null int64
job                41188 non-null object
marital            41188 non-null object
education          41188 non-null object
default            41188 non-null object
housing            41188 non-null object
loan               41188 non-null object
contact            41188 non-null object
month              41188 non-null object
day_of_week        41188 non-null object
duration           41188 non-null int64
campaign           41188 non-null int64
pdays             41188 non-null int64
previous           41188 non-null int64
poutcome           41188 non-null object
emp.var.rate       41188 non-null float64
cons.price.idx     41188 non-null float64
cons.conf.idx      41188 non-null float64
euribor3m          41188 non-null float64
nr.employed        41188 non-null float64
dtypes: float64(5), int64(5), object(10)
memory usage: 6.3+ MB
None
```

- All the features are having 41188 non-null entries and hence there are no null values
- There are 10 continuous and 10 categorical attributes

---

**Make the categorical variables of dtype 'category' and label encode**

```
In [84]: for col in df_x.select_dtypes(include=['object']):
          # Changing the dtype of categorical columns to 'category' reduces memory usage
          df_x[col] = df_x[col].astype('category')

          categorical_attrs = list(df_x.select_dtypes(include=['category']).copy().columns)

          df_x[categorical_attrs] = df_x[categorical_attrs].apply(lambda x: x.cat.codes)

          print(getCategories('education'))
          df_x.head(2)
```

```
education categories
1 basic.4y
2 basic.6y
3 basic.9y
4 high.school
5 illiterate
6 professional.course
7 university.degree
8 unknown
None
```

```
Out [84]:
```

	age	job	marital	education	default	housing	loan	contact	month	\
0	56	3	1	0	0	0	0	1	6	
1	57	7	1	3	1	0	0	1	6	

	day_of_week	duration	campaign	pdays	previous	poutcome	emp.var.rate	\
0		1	261	1	999	0	1	1.1
1		1	149	1	999	0	1	1.1

	cons.price.idx	cons.conf.idx	euribor3m	nr.employed
0	93.994	-36.4	4.857	5191.0
1	93.994	-36.4	4.857	5191.0

```
In [85]: print(df_x.info(verbose=False))
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 41188 entries, 0 to 41187
Columns: 20 entries, age to nr.employed
dtypes: float64(5), int64(5), int8(10)
memory usage: 3.5 MB
None
```

Memory usage decreased from 6.3+MB to 3.5MB by making the categorical attributes datatype as **category**

```
In [86]: labels = ["Opened term subscription", "Didn't open term subscription"]
         df_y['y'].value_counts().plot.pie(labels=labels, explode=[0, 0.1]);
```



There are more records with the output of not opening a term subscription

```
In [87]: print(df_x.describe())
```

	age	job	marital	education	default \
count	41188.00000	41188.00000	41188.00000	41188.00000	41188.00000
mean	40.02406	3.72458	1.172769	3.747184	0.208872
std	10.42125	3.59456	0.608902	2.136482	0.406686
min	17.00000	0.00000	0.000000	0.000000	0.000000
25%	32.00000	0.00000	1.000000	2.000000	0.000000
50%	38.00000	2.00000	1.000000	3.000000	0.000000
75%	47.00000	7.00000	2.000000	6.000000	0.000000
max	98.00000	11.00000	3.000000	7.000000	2.000000

	housing	loan	contact	month	day_of_week \
count	41188.00000	41188.00000	41188.00000	41188.00000	41188.00000
mean	1.071720	0.327425	0.365252	4.230868	2.004613
std	0.985314	0.723616	0.481507	2.320025	1.397575
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	3.000000	1.000000
50%	2.000000	0.000000	0.000000	4.000000	2.000000
75%	2.000000	0.000000	1.000000	6.000000	3.000000
max	2.000000	2.000000	1.000000	9.000000	4.000000

	duration	campaign	pdays	previous	poutcome \
count	41188.00000	41188.00000	41188.00000	41188.00000	41188.00000
mean	258.285010	2.567593	962.475454	0.172963	0.930101
std	259.279249	2.770014	186.910907	0.494901	0.362886
min	0.000000	1.000000	0.000000	0.000000	0.000000
25%	102.000000	1.000000	999.000000	0.000000	1.000000

50%	180.000000	2.000000	999.000000	0.000000	1.000000
75%	319.000000	3.000000	999.000000	0.000000	1.000000
max	4918.000000	56.000000	999.000000	7.000000	2.000000

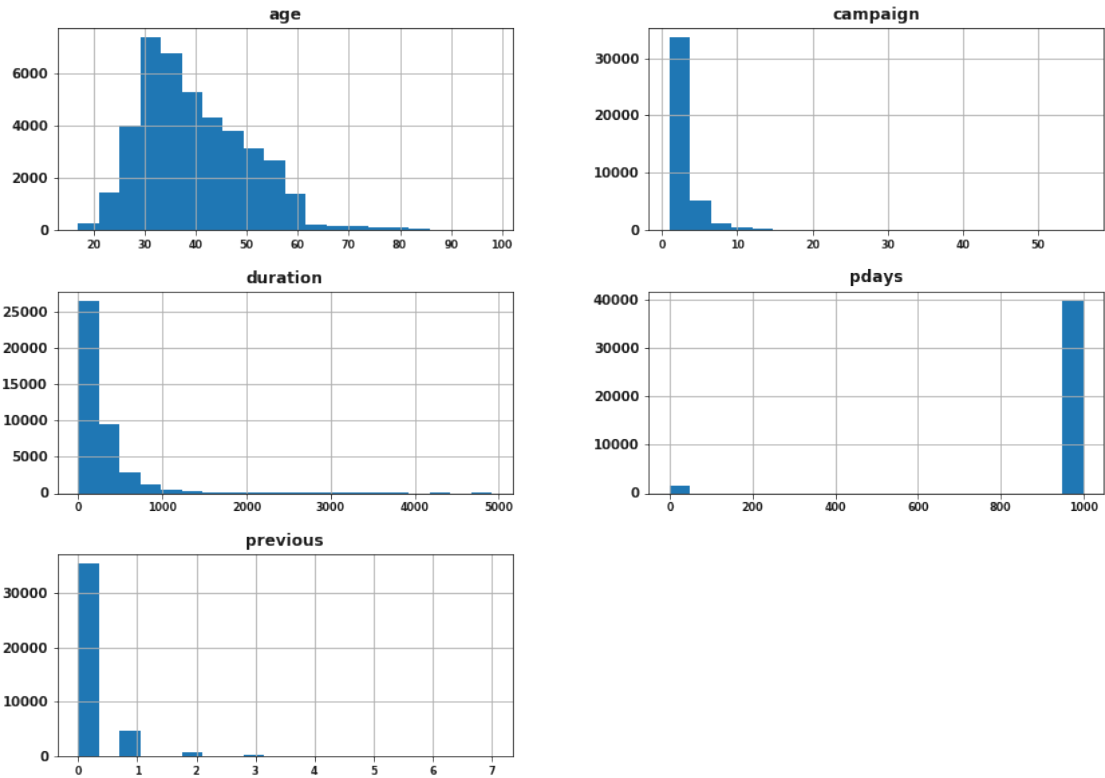
	emp.var.rate	cons.price.idx	cons.conf.idx	euribor3m	nr.employed
count	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000
mean	0.081886	93.575664	-40.502600	3.621291	5167.035911
std	1.570960	0.578840	4.628198	1.734447	72.251528
min	-3.400000	92.201000	-50.800000	0.634000	4963.600000
25%	-1.800000	93.075000	-42.700000	1.344000	5099.100000
50%	1.100000	93.749000	-41.800000	4.857000	5191.000000
75%	1.400000	93.994000	-36.400000	4.961000	5228.100000
max	1.400000	94.767000	-26.900000	5.045000	5228.100000

- Average age of the people in the dataset is 40 with std of 10.42
- Min. age is 17
- Max. age is 98
- 75% of the people have are aged below 47
- $47 * 1.5 = 70 \Rightarrow$  ages above 70 are outliers
- Average duration of the people speaking in the dataset is 258s with std of 259s
- Min. duration is 0s
- Max. duration is 4918s
- 75% of the people spoke for less than 319s
- $319 * 1.5 = 478.5 \Rightarrow$  above 478.5s are outliers

---

```
In [88]: df_num = df_x.select_dtypes(include=['int64']).copy()

df_num.hist(figsize=(14, 10), bins=20, xlabelsize=8);
```




Numerical columns except **age** and **day** are positively skewed  
**duration** is dropped as it is highly correlated to the outcome as suggested in the dataset description.

```
In [89]: df_num = df_x.select_dtypes(include=['int64', 'float64']).copy()

corr = df_num.corr()

plt.figure(figsize=(10, 10))
sns.heatmap(corr,
            cmap='viridis', vmax=1.0, vmin=-1.0, linewidths=0.1,
            annot=True, annot_kws={"size": 8}, square=True);
```






- **emp.var.rate** , **euribor3m** and **nr.employed** columns are highly correlated
- **pdays** and **previous** columns are negatively correlated

---

### Visualising the distribution of the categorical attributes

```
In [90]: plt.figure(figsize=(30, 30))

for index, col in enumerate(df_x[categorical_attrs]):
    plt.subplot(5, 2, index+1)
    df[col].value_counts().plot.pie()
```




```
In [91]: df_x.drop(['duration'], axis=1);
```

```
In [92]: plt.figure(figsize=(20, 6));
```

```
plt.subplot(1, 2, 1)
df[df['y'] == 'yes']['day_of_week'].value_counts().plot.pie()

plt.subplot(1, 2, 2)
df[df['y'] == 'no']['day_of_week'].value_counts().plot.pie()
```

```
Out[92]: <matplotlib.axes._subplots.AxesSubplot at 0x1b4d529cc88>
```



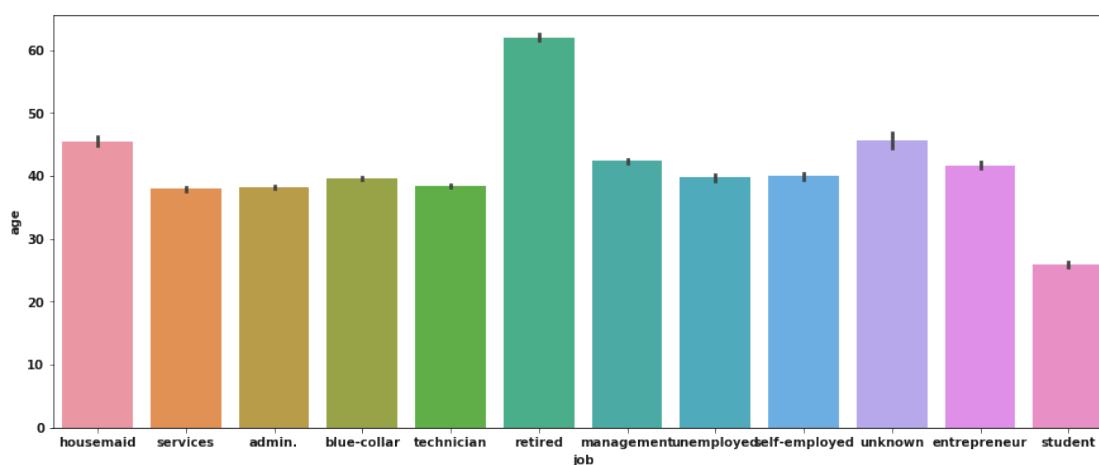
The **day\_of\_week** attributes are equally distributed and hence won't aid in classification

```
In [93]: df_x.drop(['day_of_week'], axis=1);
```

## Age


```
In [94]: plt.figure(figsize=(15, 6))
        '''
        for index, col in enumerate(df_cat.columns):
            plt.subplot(5, 2, index+1)
            #plt.subplots_adjust(hspace=1)
            sns.barplot(x=col, y='age', hue='y', data=df)#, estimator=lambda x: len(x) / len(
            #if(index==3): break
        '''
        sns.barplot(x='job', y='age', data=df)
```

```
Out[94]: <matplotlib.axes._subplots.AxesSubplot at 0x1b4d844fa20>
```



Retired people are having higher median age  
**Job**

```
In [95]: df_x['job'].value_counts().plot.pie(autopct='%d%%', figsize=(10,5));
```



Blue collar and management jobs are the occupations that are majority

## Preprocessing

```
In [96]: from sklearn.preprocessing import StandardScaler
```

```
X = df_x.values
y = df_y.values
```

```
scaler = StandardScaler()
stdScaledX = scaler.fit_transform(X)
# summarize transformed data
np.set_printoptions(precision=3)
print(stdScaledX[0:2,:])
```

```
[[ 1.533 -0.202 -0.284 -1.754 -0.514 -1.088 -0.452  1.318  0.763 -0.719
  0.01  -0.566  0.195 -0.349  0.193  0.648  0.723  0.886  0.712  0.332]
 [ 1.629  0.911 -0.284 -0.35  1.945 -1.088 -0.452  1.318  0.763 -0.719
 -0.422 -0.566  0.195 -0.349  0.193  0.648  0.723  0.886  0.712  0.332]]
```

## Checking class imbalance

```

In [97]: # number of samples who took term subscription
num_yes = len(df_y[df_y['y']==1])

# indices of samples that didnt take term subscription
no_indices = df_y[df_y['y'] == 0].index

# Random sample non term subscription
random_indices = np.random.choice(no_indices,num_yes, replace=False)

# Find the indices of term subscription
yes_indices = df_y[df_y['y']==1].index


# Concat yes indices with sample non ones
under_sample_indices = np.concatenate([yes_indices,random_indices])

# Get Balance Dataframe
under_sample = df.loc[under_sample_indices]

under_sample['y'].value_counts().plot.pie()

```

Out[97]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1b4d64fab38>



```

In [98]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)

```

### 1.0.1 Find optimal parameters for the classifiers

#### 1. Logistic Regression

```
In [99]: parameters = {'class_weight':['balanced', None]}

clf =GridSearchCV(LogisticRegression(random_state=42), parameters)
clf.fit(X_train, y_train)

print("Best parameters are: {}".format(clf.best_params_))
print(clf.cv_results_['mean_fit_time'])
print(clf.cv_results_['mean_test_score'])

Best parameters are: {'class_weight': None}
[1.005 0.382]
[0.854 0.908]
```

---

## 2. SVM

```
In [100]: parameters = {'C':[1, 5, 10], 'class_weight':['balanced', None]}

clf =GridSearchCV(LinearSVC(random_state=42), parameters)
clf.fit(X_train, y_train)

print("Best parameters are: {}".format(clf.best_params_))

Best parameters are: {'C': 1, 'class_weight': 'balanced'}
```

---

## Multi Layer Perceptron

```
In [101]: parameters = {'activation':['tanh', 'relu'], 'hidden_layer_sizes':[(100,) , (100, 100)]}

clf =GridSearchCV(MLPClassifier(random_state=42), parameters)
clf.fit(X_train, y_train)

print("Best parameters are: {}".format(clf.best_params_))

Best parameters are: {'activation': 'tanh', 'hidden_layer_sizes': (100, 100)}
```

---



---

## Classify with un-normalised features

```
In [102]: classifiers = {'Logistic Regression': LogisticRegression(random_state=42),
                        'Multilayer Perceptron': MLPClassifier(activation='tanh', hidden_layer_sizes=(100, 100)),
                        'Support Vector Classifier': LinearSVC(random_state=42)}

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)

classify_and_compare(classifiers, X_train, y_train, X_test, y_test)
```

```
Logistic Regression
      precision    recall  f1-score   support

0         0.93      0.98      0.95      9144
1         0.68      0.42      0.52      1153

avg / total         0.90      0.91      0.90     10297
```

```
Multilayer Perceptron
      precision    recall  f1-score   support

0         0.93      0.97      0.95      9144
1         0.62      0.43      0.51      1153

avg / total         0.90      0.91      0.90     10297
```

```
Support Vector Classifier
      precision    recall  f1-score   support

0         0.93      0.97      0.95      9144
1         0.64      0.47      0.54      1153


avg / total         0.90      0.91      0.90     10297
```

```

      Classifier  Accuracy  Precision Score  Recall Score \
0      Logistic Regression  0.912402      0.803615      0.695252
0      Multilayer Perceptron  0.907255      0.777226      0.700691
0      Support Vector Classifier  0.910654      0.786573      0.716627
```

```

      F1-Score  roc-auc_Score
0  0.733453      0.695252
0  0.730383      0.700691
0  0.744805      0.716627
```




---

Using standard scaled and undersampled dataset from now on

```
In [103]: normX = stdScaledX[under_sample_indices]
          normy = y[under_sample_indices]
```

### Classification with normalised features

```
In [104]: classifiers = {'Logistic Regression': LogisticRegression(random_state=42),
                        'Multilayer Perceptron': MLPClassifier(activation='tanh', hidden_layer_sizes=(100, 100)),
                        'Support Vector Classifier': LinearSVC(random_state=42)}
          X_train, X_test, y_train, y_test = train_test_split(normX, normy, test_size=0.25, random_state=42)
          classify_and_compare(classifiers, X_train, y_train, X_test, y_test)
```

```
Logistic Regression
      precision    recall  f1-score   support

     0       0.87      0.84      0.85      1161
     1       0.84      0.87      0.86      1159

 avg / total       0.85      0.85      0.85      2320
```

```
Multilayer Perceptron
      precision    recall  f1-score   support
```



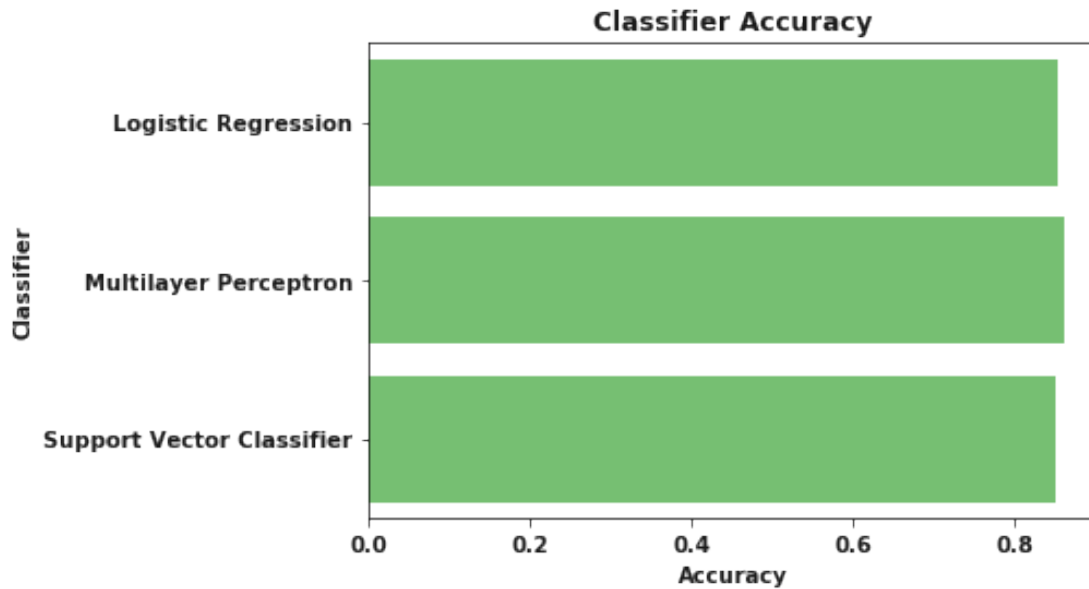
0	0.87	0.85	0.86	1161
1	0.85	0.88	0.86	1159
avg / total	0.86	0.86	0.86	2320

#### Support Vector Classifier

	precision	recall	f1-score	support
0	0.86	0.84	0.85	1161
1	0.84	0.86	0.85	1159
avg / total	0.85	0.85	0.85	2320

	Classifier	Accuracy	Precision Score	Recall Score	\
0	Logistic Regression	0.853879	0.854234	0.853893	
0	Multilayer Perceptron	0.861638	0.861962	0.861651	
0	Support Vector Classifier	0.850862	0.851124	0.850874	

	F1-Score	roc-auc_Score
0	0.853846	0.853893
0	0.861610	0.861651
0	0.850837	0.850874




---

**Classification with polynomial features**

```

In [105]: from sklearn.preprocessing import PolynomialFeatures
poly = PolynomialFeatures()

X_train, X_test, y_train, y_test = train_test_split(normX, normy, test_size=0.25, ra

X_train = poly.fit_transform(X_train)
#print(len(poly.get_feature_names()))
#print(poly.get_params())

X_test = poly.transform(X_test)

classifiers = {'Logistic Regression':LogisticRegression(random_state=42),
               }

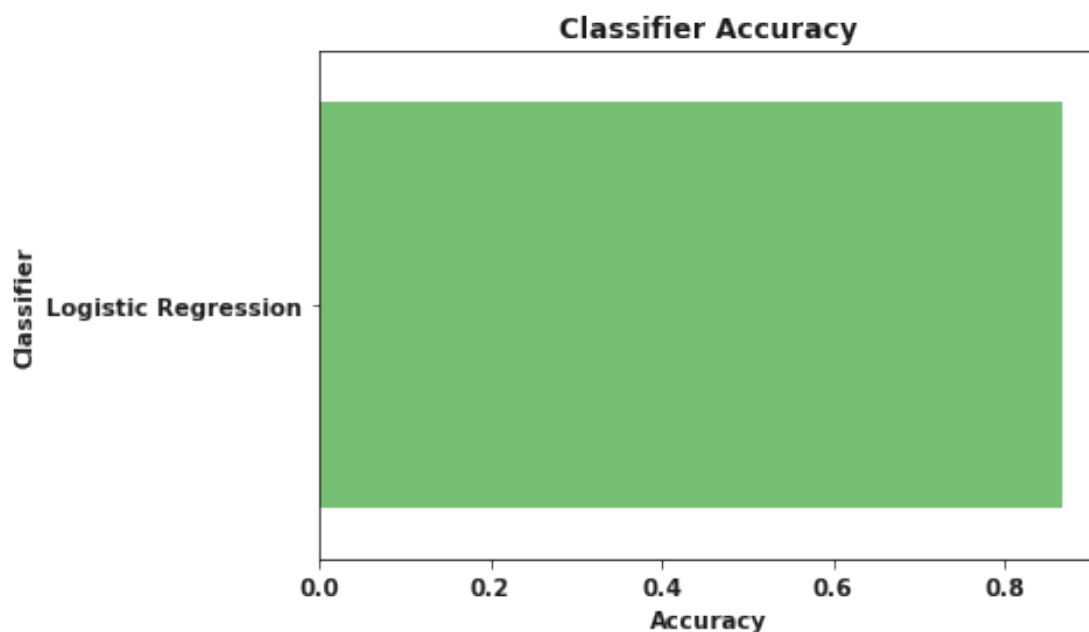
classify_and_compare(classifiers, X_train, y_train, X_test, y_test)

```

Logistic Regression				
	precision	recall	f1-score	support
0	0.89	0.84	0.86	1161
1	0.85	0.89	0.87	1159
avg / total	0.87	0.87	0.87	2320

	Classifier	Accuracy	Precision Score	Recall Score	F1-Score \
0	Logistic Regression	0.86681	0.867627	0.866831	0.866741

	roc-auc_Score
0	0.866831



## 1.0.2 PCA

```
In [106]: from sklearn.decomposition import PCA
```

```
X_train, X_test, y_train, y_test = train_test_split(normX, normy, test_size=0.25, random_state=42)

# Make an instance of the Model
# 80% variance is retained
pca = PCA(.80)
pca.fit(X_train)

X_train = pca.transform(X_train)
X_test = pca.transform(X_test)

print(X_train.shape)
```

```
(6960, 10)
```

### Classification after dimensionality reduction using PCA

```
In [107]: classifiers = {'Logistic Regression': LogisticRegression(random_state=42),
                        'Multilayer Perceptron': MLPClassifier(activation='tanh', hidden_layer_sizes=(100, 100)),
                        'Support Vector Classifier': LinearSVC(random_state=42)}

# Train the classifiers
for name, clf in classifiers.items():
    clf.fit(X_train, y_train)
```

```
classify_and_compare(classifiers, X_train, y_train, X_test, y_test)
```

#### Logistic Regression

	precision	recall	f1-score	support
0	0.84	0.84	0.84	1161
1	0.84	0.84	0.84	1159
avg / total	0.84	0.84	0.84	2320

#### Multilayer Perceptron

	precision	recall	f1-score	support
0	0.88	0.83	0.86	1161
1	0.84	0.89	0.87	1159
avg / total	0.86	0.86	0.86	2320

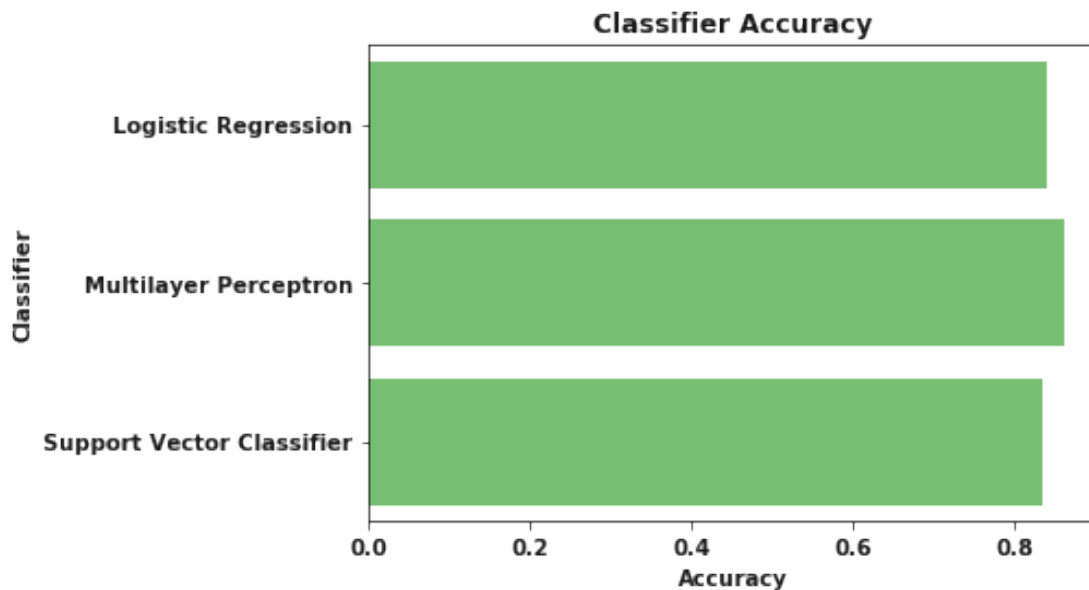
#### Support Vector Classifier

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.83	0.84	0.84	1161
1	0.84	0.83	0.83	1159
avg / total	0.84	0.84	0.84	2320

	Classifier	Accuracy	Precision Score	Recall Score	\
0	Logistic Regression	0.840948	0.840955	0.840946	
0	Multilayer Perceptron	0.862069	0.863164	0.862093	
0	Support Vector Classifier	0.835345	0.835382	0.835340	

	F1-Score	roc-auc_Score
0	0.840947	0.840946
0	0.861970	0.862093
0	0.835339	0.835340



### PCA with polynomial features

```
In [108]: X_train, X_test, y_train, y_test = train_test_split(normX, normy, test_size=0.25, random_state=42)

X_train = poly.fit_transform(X_train)
X_test = poly.transform(X_test)

print(len(poly.get_feature_names()))
```

```

pca = PCA(.80)
pca.fit(X_train)

X_train = pca.transform(X_train)
X_test = pca.transform(X_test)

print(X_train.shape)

classifiers = {'Logistic Regression': LogisticRegression(random_state=42),
               'Multilayer Perceptron': MLPClassifier(activation='tanh', hidden_layer_sizes=(100, 100)),
               'Support Vector Classifier': LinearSVC(random_state=42)}

classify_and_compare(classifiers, X_train, y_train, X_test, y_test)

```

231  
(6960, 33)

Logistic Regression

	precision	recall	f1-score	support
0	0.80	0.83	0.81	1161
1	0.82	0.79	0.80	1159
avg / total	0.81	0.81	0.81	2320

Multilayer Perceptron

	precision	recall	f1-score	support
0	0.84	0.84	0.84	1161
1	0.84	0.84	0.84	1159
avg / total	0.84	0.84	0.84	2320

Support Vector Classifier

	precision	recall	f1-score	support
0	0.79	0.84	0.81	1161
1	0.83	0.77	0.80	1159
avg / total	0.81	0.81	0.81	2320

	Classifier	Accuracy	Precision Score	Recall Score	\
0	Logistic Regression	0.809052	0.809753	0.809031	
0	Multilayer Perceptron	0.837931	0.837931	0.837931	
0	Support Vector Classifier	0.805603	0.806758	0.805577	

	F1-Score	roc-auc_Score
0	0.808936	0.809031

0	0.837931	0.837931
0	0.805411	0.805577

